

DATA 605 : Final Exam - Problem2

Ramnivas Singh

12/14/2021

Problem 2 - Digit Recognizer

Digit Recognizer is one of the basic and first problem that a budding Machine Learning engineer should try their hands on. It is a simple problem where the challenge is to recognize hand-written digits.

1. Go to Kaggle.com and build an account if you do not already have one. It is free.

Answer : An existing account used to access Kaggle.com

2. Go to <https://www.kaggle.com/c/digit-recognizer/overview>, accept the rules of the competition, and download the data. You will not be required to submit work to Kaggle, but you do need the data.

Answer :

3. Using the training.csv file, plot representations of the first 10 images to understand the data format. Go ahead and divide all pixels by 255 to produce values between 0 and 1. (This is equivalent to min-max scaling.)

```
train <- read.csv("train.csv")
test <- read.csv("test.csv")
# Total Rows of training dataset
nrow(train) # Dataset has 4200 records
```

Answer :

```
## [1] 42000
```

```
# Total columns of training dataset
ncol(train)
```

```
## [1] 785
```

```
# Total Rows of training dataset
nrow(test) # Dataset has 4200 records
```

```
## [1] 28000
```

```
# Total columns of training dataset
ncol(test)
```

```
## [1] 784
```

```
# Print top records from train dataset
head(train[1:10])
```

```
##   label pixel0 pixel1 pixel2 pixel3 pixel4 pixel5 pixel6 pixel7 pixel8
## 1     1      0      0      0      0      0      0      0      0      0
## 2     0      0      0      0      0      0      0      0      0      0
## 3     1      0      0      0      0      0      0      0      0      0
## 4     4      0      0      0      0      0      0      0      0      0
## 5     0      0      0      0      0      0      0      0      0      0
## 6     0      0      0      0      0      0      0      0      0      0
```


```
summary(train[train$label==1, 408])
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      0.0   253.0   253.0   246.5   254.0   255.0
```

```
summary(train[train$label==0, 408])
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      0.000   0.000   0.000   4.517   0.000  255.000
```

```
m<-matrix(unlist(train[12,-1]), nrow=28, byrow = T)
image(m, col = grey.colors(255))
```



```
flip <- function(matrix){
  apply(matrix, 2, rev)
}

digit<-function(x){
  m<-matrix(unlist(x), nrow=28, byrow=T)
  m<-t(apply(m, 2, rev))
  image(m, col=grey.colors(255))
}

par(mfrow=c(3,4))

for(i in 1:10){
  digit(train[i, -1])
}

# divide all pixels by 255 to produce values between 0 and 1.


train_255 <- train/255.0
test_255  <- test/255.0
head(train_255[1:10])
```

```
##          label pixel0 pixel1 pixel2 pixel3 pixel4 pixel5 pixel6 pixel7 pixel8
## 1 0.003921569      0      0      0      0      0      0      0      0      0
## 2 0.000000000      0      0      0      0      0      0      0      0      0
```

```
## 3 0.003921569      0      0      0      0      0      0      0      0      0
## 4 0.015686275      0      0      0      0      0      0      0      0      0
## 5 0.000000000      0      0      0      0      0      0      0      0      0
## 6 0.000000000      0      0      0      0      0      0      0      0      0
```

```
# In dataset-remove the first column of label. Create a new data set for this operation apply Min-Max n
normalize <- function(x, na.rm = TRUE) {
  return((x- min(x)) /(max(x)-min(x)))
}

train_b <- train %>% select( 2:ncol(.) )
train_b<- as.data.frame(lapply(train_b[, -1], normalize))
par(mfrow=c(3,4))
```



```
#Removing Missing values from training(train) dataframe
train[is.na(train)] <- 0
```

4. What is the frequency distribution of the numbers in the dataset?


Answer : Frequency Distribution : A frequency distribution is a representation, either in a graphical or tabular format, that displays the number of observations within a given interval or categories. It is also called the Frequency Distribution table. Given this dataset, frequency distribution table shows occurrence

of various labels in training dataset. To investigate the balance of the data for each label, function will plot all of them along with their name. A plot for the same is below

```
# Table shows numeric representation of frequency distribution  
table(train$label)
```

```
##  
##    0    1    2    3    4    5    6    7    8    9  
## 4132 4684 4177 4351 4072 3795 4137 4401 4063 4188
```

```
# Bar plot chart to show frequency distribution  
barplot(table(train$label), main="Total Number of Digits (Training dataset)", col=brewer.pal(10,"Set1")  
        xlab="Numbers", ylab = "Frequency of Numbers")
```

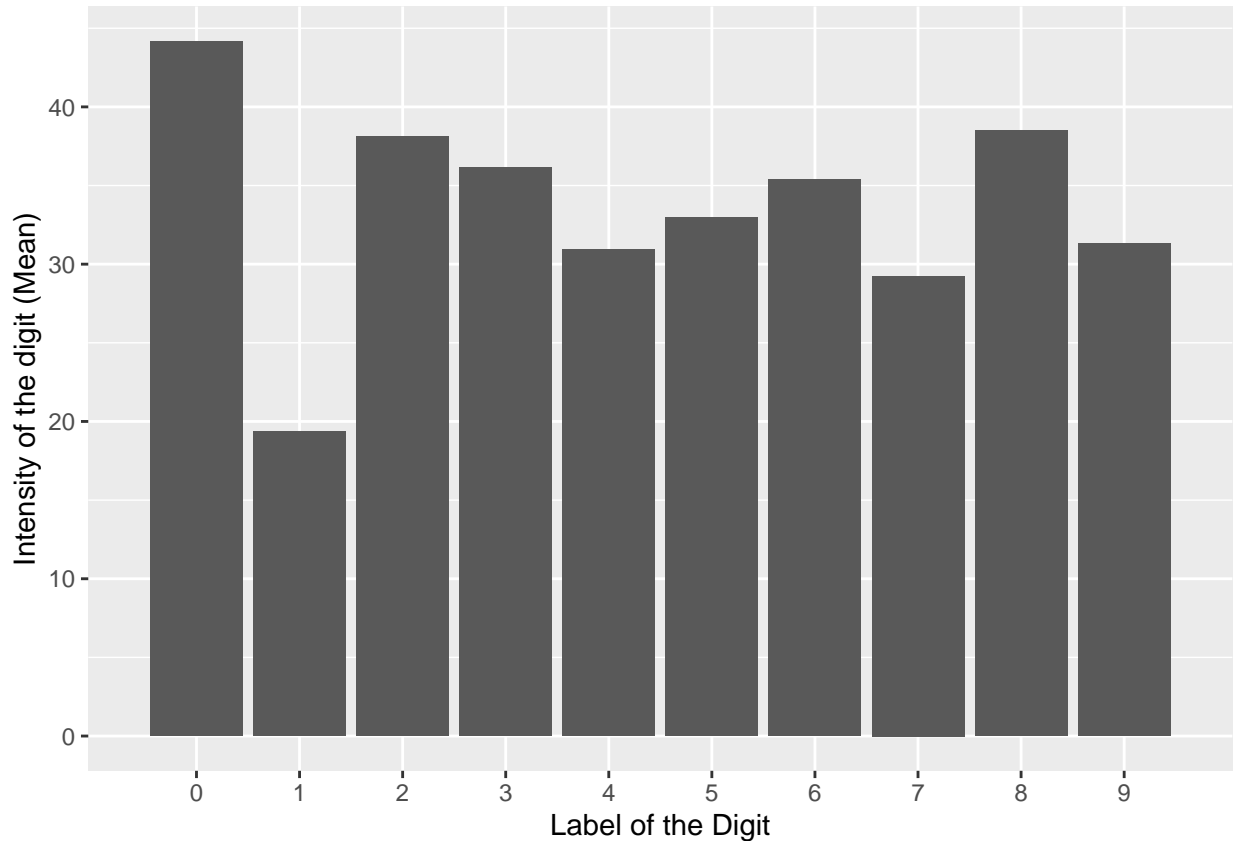


5. For each number, provide the mean pixel intensity. What does this tell you?

Answer : To find mean pixel intensity of all of the pixels, we can reshape the dataframe. The output is actually a series, indexed by the original index as well as the pixel label:

```
# Lets calculate mean of each row in training data set to find pixel intensity of all of the pixels  
train$intensity <- apply(train[,-1], 1, mean)  
intensity_by_label <- aggregate (train$intensity, FUN = mean, by = list(train$label))
```


```
# Now lets plot intensity_by_label
ggplot(data=intensity_by_label, aes(x=Group.1, y = x)) +
  geom_bar(stat="identity")+scale_x_discrete(limits=0:9) + xlab("Label of the Digit") +
  ylab("Intensity of the digit (Mean)")
```



We see that pixel values have different intensity across the dataset. 0, 8, 2, 6 are with higher mean pixel intensity where as 1, 4, 7, and with low mean pixel intensity. If we want to bin the intensity values into statistical quantiles, we can do that. Overall digit “0” is the most intense and “1” is the less intense.


Lets plot Histogram for 0, 8, 2, 6

```
grid.arrange(qplot(subset(train, label ==0)$intensity, binwidth = .5, xlab = "Histogram for 0"),
qplot(subset(train, label ==8)$intensity, binwidth = .5, xlab = "Histogram for 8"),
qplot(subset(train, label ==2)$intensity, binwidth = .5, xlab = "Histogram for 2"),
qplot(subset(train, label ==6)$intensity, binwidth = .5, xlab = "Histogram for 6"),ncol = 4)
```



Lets plot Histogram for 1, 4, 7, 9

```
grid.arrange(qplot(subset(train, label ==0)$intensity, binwidth = .5, xlab = "Histogram for 1"),
qplot(subset(train, label ==8)$intensity, binwidth = .5, xlab = "Histogram for 4"),
qplot(subset(train, label ==2)$intensity, binwidth = .5, xlab = "Histogram for 7"),
qplot(subset(train, label ==6)$intensity, binwidth = .5, xlab = "Histogram for 9"),ncol = 4)
```




6. Reduce the data by using principal components that account for 95% of the variance. How many components did you generate? Use PCA to generate all possible components (100% of the variance). How many components are possible? Why?

```
pcaCharts <- function(x) {
  x.var <- x$sdev ^ 2
  x.pvar <- x.var/sum(x.var)
  par(mfrow=c(1,1))
  plot(x.pvar,xlab="Principal component", ylab="Proportion of variance explained", ylim=c(0,1), type="b")
  plot(cumsum(x.pvar),xlab="Principal component", ylab="Cumulative Proportion of variance explained", type="b")
  screeplot(x,type="l")
  par(mfrow=c(1,1))
}

#Reducing data using PCA
train_norm<-as.matrix(train[,1:255])/255
train_norm_cov <- cov(train_norm)
pca <- prcomp(train_norm_cov)

pcaCharts(pca)
```

```
# Calculate the variance explained by each principal component
variance_explained<-as.data.frame(pca$sdev^2/sum(pca$sdev^2))
variance_explained<-cbind(1:785, cumsum(variance_explained))
colnames(variance_explained)<-c("Number", "Variance")
variance_explained<-as.data.frame(variance_explained)
head(variance_explained,100)
```

```
##      Number  Variance
## 1         1 0.2533019
## 2         2 0.4211143
## 3         3 0.5443375
## 4         4 0.6382797
## 5         5 0.7099622
## 6         6 0.7691849
## 7         7 0.8028326
## 8         8 0.8302254
## 9         9 0.8531247
## 10        10 0.8711668
## 11        11 0.8853364
## 12        12 0.8986528
## 13        13 0.9080900
## 14        14 0.9174678
## 15        15 0.9256073
## 16        16 0.9328088
## 17        17 0.9384871
## 18        18 0.9438281
```

## 19	19 0.9483989
## 20	20 0.9527310
## 21	21 0.9564960
## 22	22 0.9598711
## 23	23 0.9629199
## 24	24 0.9656478
## 25	25 0.9682119
## 26	26 0.9705036
## 27	27 0.9726586
## 28	28 0.9746368
## 29	29 0.9764279
## 30	30 0.9779677
## 31	31 0.9793804
## 32	32 0.9807166
## 33	33 0.9818930
## 34	34 0.9830205
## 35	35 0.9840633
## 36	36 0.9850209
## 37	37 0.9858696
## 38	38 0.9866471
## 39	39 0.9873878
## 40	40 0.9880989
## 41	41 0.9887694
## 42	42 0.9894178
## 43	43 0.9899900
## 44	44 0.9905075
## 45	45 0.9909903
## 46	46 0.9914504
## 47	47 0.9918771
## 48	48 0.9922717
## 49	49 0.9926425
## 50	50 0.9929780
## 51	51 0.9933007
## 52	52 0.9936134
## 53	53 0.9938955
## 54	54 0.9941623
## 55	55 0.9944194
## 56	56 0.9946575
## 57	57 0.9948887
## 58	58 0.9951032
## 59	59 0.9953134
## 60	60 0.9955119
## 61	61 0.9956998
## 62	62 0.9958863
## 63	63 0.9960555
## 64	64 0.9962159
## 65	65 0.9963657
## 66	66 0.9965039
## 67	67 0.9966386
## 68	68 0.9967638
## 69	69 0.9968865
## 70	70 0.9970029
## 71	71 0.9971165
## 72	72 0.9972235


```
## 73      73 0.9973257
## 74      74 0.9974232
## 75      75 0.9975134
## 76      76 0.9976006
## 77      77 0.9976845
## 78      78 0.9977623
## 79      79 0.9978322
## 80      80 0.9978985
## 81      81 0.9979637
## 82      82 0.9980280
## 83      83 0.9980911
## 84      84 0.9981509
## 85      85 0.9982082
## 86      86 0.9982642
## 87      87 0.9983192
## 88      88 0.9983697
## 89      89 0.9984186
## 90      90 0.9984655
## 91      91 0.9985092
## 92      92 0.9985516
## 93      93 0.9985929
## 94      94 0.9986328
## 95      95 0.9986710
## 96      96 0.9987087
## 97      97 0.9987440
## 98      98 0.9987787
## 99      99 0.9988124
## 100     100 0.9988451
```

There are around 20 components generated at 95% of variance


7. Plot the first 10 images generated by PCA. They will appear to be noise. Why?

Answer : Because PCs of higher eigenvalues to show more generic features. The specialized features are also being added for more PCs. So removing removing some PCs with lower eigenvalues actually acting as some sort of regularization and your model is only learning the more general features. If you take all of them the 100% of the data-variations will be restored like the original dimensions.

```
labelClasses <- factor(train$label)
plot(main="",pca$x, col = labelClasses)
```



```
for(i in 1:10){  
  digit(pca$x[i, -1])  
}
```




8. Now, select only those images that have labels that are 8's. Re-run PCA that accounts for all of the variance (100%). Plot the first 10 images. What do you see?


```
train_sub_8<-subset(train, label ==8)
nrow(train_sub_8)
```


Answer :

```
## [1] 4063
```


```
pca <- prcomp(train_sub_8)
pcaCharts(pca)
```







```
par(mfrow=c(3,4))
for(i in 1:10){
  digit(train_sub_8[i, -1])
}
```



In above images, different type of 8 letters seems to be visible. ***

9. An incorrect approach to predicting the images would be to build a linear regression model with y as the digit values and X as the pixel matrix. Instead, we can build a multinomial model that classifies the digits. Build a multinomial model on the entirety of the training set. Then provide its classification accuracy (percent correctly identified) as well as a matrix of observed versus forecast values (confusion matrix). This matrix will be a 10 x 10, and correct classifications will

be on the diagonal.

NOTE: Please note models used below for this problem are running on small dataset extracted from train.csv file. Normal laptop doesnt return with the result even after many hours of continuous processing.

Answer : To build a multinomial model on the entirety of the training, I started with training dataset by splitting dataset further into to train and test where 70% data used for training the model and 30% dataset for testing it. Approach is run multinom function which will than return multinomModel than use multinomModel in predict function for probable results. Running this model produced around 93% accuracy.

Lines below for this model are commented, this model for 42K records takes around 4 days of time to return the results.

```
train <- read.csv("train_1.csv")

train[is.na(train)] <- 0
sample_size = round(nrow(train)*.70) # setting what is 70%
```

```

index <- sample(seq_len(nrow(train)), size = sample_size)

train <- train[index, ]
test <- train[-index, ]

#Build Multinomial Model
#multinomModel <- multinom(label ~., family = "multinomial", data = train, MaxNWts=100000, maxit=10);
#summary (multinomModel) # model summary

#predicted_scores <- predict (multinomModel, test, "probs") # predict on multinomModel data
#predicted_class <- predict (multinomModel, test) # model summary

#Confusion Matrix and Misclassification Error
#table(predicted_class, test$label)
#mean(as.character(predicted_class) != as.character(test$label))

```

A classification accuracy of 93.3% from multinom model is probably too high. We should try other ML approaches as well for this problem.

*** Gradient boosted trees model for multinomial *** : Gradient boosting is a machine learning technique used in regression and classification tasks, among others. It gives a prediction model in the form of an ensemble of weak prediction models, which are typically decision trees. Gradient boosted trees also run directly on the multiclass labels. The model performs much better if I increase the interaction depth slightly. Increasing it past 2-3 is beneficial in large models, but rarely useful with smaller cases like this. I could also play with the learning rate, but won't fiddle with that here for now.

```

# Model fitting
Xtrain <- as.matrix(train)
Xtest <- as.matrix(test)
ytrain <- train[,1]
ytest <- test[,1]
gbm_result <- gbm.fit(Xtrain, factor(ytrain), distribution="multinomial", n.trees=500, interaction.dep

```

## Iter	TrainDeviance	ValidDeviance	StepSize	Improve
## 1	2.3026	nan	0.0010	0.0161
## 2	2.2939	nan	0.0010	0.0147
## 3	2.2860	nan	0.0010	0.0161
## 4	2.2774	nan	0.0010	0.0120
## 5	2.2704	nan	0.0010	0.0126
## 6	2.2633	nan	0.0010	0.0141
## 7	2.2557	nan	0.0010	0.0131
## 8	2.2487	nan	0.0010	0.0142
## 9	2.2411	nan	0.0010	0.0139
## 10	2.2336	nan	0.0010	0.0134
## 20	2.1644	nan	0.0010	0.0117
## 40	2.0321	nan	0.0010	0.0100
## 60	1.9249	nan	0.0010	0.0077
## 80	1.8270	nan	0.0010	0.0076
## 100	1.7395	nan	0.0010	0.0088
## 120	1.6551	nan	0.0010	0.0079
## 140	1.5792	nan	0.0010	0.0065
## 160	1.5100	nan	0.0010	0.0058
## 180	1.4455	nan	0.0010	0.0073

```
##      200      1.3852      nan      0.0010      0.0057
##      220      1.3283      nan      0.0010      0.0059
##      240      1.2753      nan      0.0010      0.0041
##      260      1.2263      nan      0.0010      0.0042
##      280      1.1778      nan      0.0010      0.0053
##      300      1.1306      nan      0.0010      0.0044
##      320      1.0856      nan      0.0010      0.0035
##      340      1.0442      nan      0.0010      0.0037
##      360      1.0028      nan      0.0010      0.0027
##      380      0.9643      nan      0.0010      0.0033
##      400      0.9281      nan      0.0010      0.0032
##      420      0.8923      nan      0.0010      0.0034
##      440      0.8576      nan      0.0010      0.0038
##      460      0.8252      nan      0.0010      0.0027
##      480      0.7931      nan      0.0010      0.0030
##      500      0.7620      nan      0.0010      0.0026
```

```
#summary(gbm_result)
gbm_prediction <- apply(predict(gbm_result, Xtest, n.trees=gbm_result$n.trees),1,which.max) - 1L
# Prediction
gbm_prediction
```

```
##      [1] 4 3 1 7 0 1 1 0 6 0 9 0 3 9 1 3 0 3 9 9 2 5 6 4 0 4 5 7 8 8 2 6 0 2 5 4 1
##      [38] 9 2 9 9 2 4 2 1 0 5 9 7 7 4 0 7 8 4 7 2 4 4 6 2 9 5 6 5 0 7 2 5 3 3 8 5 3
##      [75] 1 2 0 4 4 1 3 8 2 0 9 7 2 2 1 6 2 8 1 0 6 8 8 5 1 0 7 7 3 2 8 4 9 4 4 5 3
##     [112] 9 5 2 5 1 1 0 4 6 0 7 3 5 5 7 4 5 1 9 1 9 9 3 2 5 4 4 0 4 6 2 4 6 1 9 5 2
##     [149] 6 5 9 6 6 0 1 5 6 7 2 3 9 7 3 9 0 7 2 6 1 2 1 1 5 3 5 6 3 8 4 1 9 5 2 2 5
##     [186] 4 3 3 1 2 0 3 6 2 5 5 1 3 4 3 9 0 7 1 7 9 3 4 8 7 0 8 1 1 8 1 7 0 6 9 6 9
##     [223] 8 6 7 4 5 5 4 8 9 6 0 7 3 1 6 4 9 5 9 1 1 9 7 2 9 1 7 7 4 6 9 0 2 9 4 2 6
##     [260] 1 2 1 7 4 1 6 5 0 0 8 9 2 6 7 0 8 9 0 7 7 9 0 5 2 3 1 3 4 5 0 0 9 9 9 3 2
##     [297] 1 7 4 2 9 1 4 6 7 6 6 4 9 5 5 0 3 2 9 8 7 4 9 6 3 2 0 2 5 1 7 2 8 8 4 1 8
##     [334] 7 6 3 7 0 8 4 3 6 4 2 9 4 0 7 2 5 1 8 2 4 4 9 6 9 8 0 8 1 5 0 9 0 8 2 3 7
##     [371] 2 9 9 1 0 0 4 6 1 8 8 2 7 5 5 8 8 8 4 7 4 8 0 7 8 8 9 7 5 6 5 5 1 3
```

Lets try Ridge regression for prediction and confusion, I'll directly use the multinomial loss function and let the R function do cross validation this time.

```
library(glmnet)
outLm <- cv.glmnet(Xtrain, ytrain, alpha=0, nfolds=3,
                  family="multinomial")
outLm
```

```
##
## Call:  cv.glmnet(x = Xtrain, y = ytrain, nfolds = 3, alpha = 0, family = "multinomial")
##
## Measure: Multinomial Deviance
##
##      Lambda Index Measure      SE Nonzero
## min 0.01624   100  0.8403 0.02915      629
## 1se 0.03753    91  0.8666 0.03246      629
```

```

predLm <- apply(predict(outLm, Xtest, s=outLm$lambda.min,
                        type="response"), 1, which.max) - 1L
predLm

```

```

## 300 1127 1520 494 1260 1001 1792 1511 790 201 1949 1661 730 1425 1756 1555
## 4 3 1 7 0 1 1 0 6 0 7 0 3 9 1 3
## 268 1536 1102 847 928 1467 1887 216 1730 310 590 243 1979 1208 683 1915
## 0 3 9 9 2 5 6 4 0 4 5 7 8 8 2 6
## 1304 1926 504 352 1341 1308 1686 1562 791 1990 50 1207 1774 1992 1739 771
## 0 2 5 4 1 9 2 9 9 2 4 2 1 0 5 9
## 49 1559 828 328 658 552 674 932 1842 883 1880 635 516 1807 1592 400
## 7 7 4 0 7 8 4 7 2 4 4 6 2 9 5 6
## 1353 1463 812 1055 1611 315 1138 1027 997 482 210 949 1097 704 951 1414
## 5 0 7 2 5 3 3 8 5 3 1 2 0 4 4 1
## 734 717 1614 205 856 1597 600 1867 13 584 398 1738 1025 1228 981 612
## 3 8 2 0 9 7 2 2 1 6 2 8 1 0 6 8
## 11 1150 425 897 1316 1066 334 1262 895 1623 678 1625 1802 298 1184 722
## 8 5 1 0 7 7 3 2 8 4 9 4 4 5 3 9
## 848 665 1679 1737 1995 1429 1711 1705 189 1093 1660 1983 1583 1152 1804 100
## 5 2 5 1 1 0 4 6 0 7 3 5 5 7 4 5
## 16 101 125 800 442 1435 1331 1113 1778 1972 920 1819 1492 560 525 1860
## 1 9 1 9 9 3 2 5 4 4 0 4 6 2 4 6
## 1993 1602 1839 25 217 390 1645 271 1241 1989 1 1833 1688 1359 1395 1626
## 1 9 5 2 6 5 9 6 6 0 1 5 6 7 2 3
## 1489 1460 1858 1258 676 505 1299 1195 709 1176 469 545 1822 1604 666 1565
## 9 7 3 9 0 7 2 6 1 2 1 1 5 3 5 6
## 229 535 483 827 1280 767 1420 1231 1439 497 66 1234 1663 1734 287 943
## 3 8 4 1 9 5 2 2 5 4 3 3 1 2 0 3
## 1728 1013 1622 1628 1436 1922 841 231 211 882 1242 1295 1849 657 388 206
## 6 2 5 5 1 3 4 3 9 0 7 1 7 9 3 4
## 1300 224 986 462 1613 1836 407 591 655 1570 290 1878 46 181 486 1229
## 8 7 0 8 1 1 8 1 7 0 6 9 6 9 8 6
## 1349 468 1766 1468 421 1795 281 160 871 157 1895 839 1309 1815 212 1827
## 7 4 5 5 4 8 9 6 0 7 3 1 6 4 9 5
## 1719 1752 1946 1746 1252 1155 1945 1480 1418 594 1510 1975 840 473 1061 953
## 9 1 1 9 7 2 9 1 7 7 4 6 9 0 2 9
## 1462 909 1921 69 279 976 238 1068 1269 128 1960 274 1286 849 1707 264
## 4 2 6 1 2 1 7 4 1 6 5 0 0 8 9 2
## 1962 1101 1765 450 304 293 680 1588 1683 975 603 1866 972 1932 1664 772
## 6 7 0 8 9 0 7 7 9 0 5 2 3 1 3 4
## 329 784 733 1110 1893 1947 773 1550 1889 885 1485 45 1259 599 1251 1714
## 5 0 0 9 9 9 3 2 1 7 4 2 9 1 4 6
## 967 161 331 1087 438 1164 461 374 285 508 979 933 337 819 1865 1283
## 7 6 5 4 9 5 5 0 3 2 9 8 7 4 9 6
## 607 1603 1077 1680 499 646 857 907 1941 172 1817 61 520 817 631 1821
## 3 2 0 2 5 1 7 2 8 8 4 1 8 7 6 3
## 609 1478 455 805 868 230 1296 474 547 1733 1169 1157 208 801 1434 1191
## 1 0 8 4 3 6 4 2 9 4 0 7 2 5 1 8
## 726 1000 44 162 1202 470 1397 1687 223 372 686 628 357 1671 1179 354
## 2 4 4 9 6 9 8 0 8 1 5 0 9 0 8 2
## 1553 1694 637 1624 54 250 672 1475 1383 899 1506 815 675 480 314 738
## 3 7 2 9 9 1 0 0 4 6 1 8 8 2 7 5
## 1372 1049 1549 1670 811 1919 79 1835 436 1011 1607 851 308 19 1984 605

```

```
##      5      8      8      8      4      7      4      8      0      7      8      8      9      7      5      6
## 175 652 970 71
##      5      5      1      3
```

```
#Mis-classification rates by class
tapply(predLm != ytest, ytest, mean)
```

```
##              0              1              2              3              4              5              6
## 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
##              7              8              9
## 0.02500000 0.00000000 0.02083333
```

```
# Confusion Metric
table(predLm,ytest)
```

```
##      ytest
## predLm 0  1  2  3  4  5  6  7  8  9
##      0 42  0  0  0  0  0  0  0  0
##      1  0 44  0  0  0  0  0  1  0
##      2  0  0 44  0  0  0  0  0  0
##      3  0  0  0 32  0  0  0  0  0
##      4  0  0  0  0 44  0  0  0  0
##      5  0  0  0  0  0 42  0  0  0
##      6  0  0  0  0  0  0 34  0  0
##      7  0  0  0  0  0  0  0 39  0
##      8  0  0  0  0  0  0  0  0 34
##      9  0  0  0  0  0  0  0  0  0 47
```

We might think that a lot 2's, 6's and 9's are being mis-classified as one another (they do look similar in some ways). Looking at the confusion matrices we see that this is not quite the case.